

# Setup

---

- install Java; check version (`java -version`)
- download Confluent platform from <https://www.confluent.io/download/>
- unzip and untar (`tar xvfz confluent-5*.tar.gz`)
- move to /opt
- create symbolic link (`ln -s confluent-5* confluent`)
- add to path (`export PATH=${PATH}:/opt/confluent/bin`)
- install the Confluent CLI (see <https://docs.confluent.io/current/cli/installing.html#scripted-installation>)
  - `curl -L https://cnfl.io/cli | sh -s -- -b /opt/confluent/bin`
- Start
  - `confluent local start ksql-server`
  - or if using Docker `docker-compose up -d`

## Confluent CLI changes from version 5.3

---

The *Confluent CLI* changed significantly in version 5.3. The most important change is the inclusion of the `local` paramter when interacting with a local development environment.

For example, to start the ksql-server

- Prior to 5.3 : `confluent start ksql-server`
- From 5.3 : `confluent local start ksql-server`

## Create a topic

---

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-factor 1 --topic USERS
```

Or you can use the more modern syntax

```
kafka-topics --bootstrap-server localhost:9092 --create --partitions 1 --replication-factor 1 --topic USERS
```

## Get started - KSQL Command Line

---

At UNIX prompt

```
kafka-console-producer --broker-list localhost:9092 --topic USERS << EOF
Alice,US
```

```
Bob,GB
Carol,AU
Dan,US
EOF
```

At KSQL prompt

```
show topics;

-- this will show nothing
print 'USERS';

print 'USERS' from beginning;

print 'USERS' from beginning limit 2;

print 'USERS' from beginning interval 2 limit 2 ;
```

## Get started - Create a stream with CSV

---

At KSQL prompt

```
create stream users_stream (name VARCHAR, countrycode VARCHAR) WITH
(KAFKA_TOPIC='USERS', VALUE_FORMAT='DELIMITED');

list streams;

-- nothing will get shown
select name, countrycode from users_stream;
```

**auto.offset.reset** - Determines what to do when there is no initial offset in Apache Kafka or if the current offset does not exist on the server. The default value in KSQL is *latest*, which means all Kafka topics are read from the latest available offset. For example, to change it to *earliest* by using the KSQL command line:

```
-- default to beginning of time
SET 'auto.offset.reset'='earliest';

-- now will see something
select name, countrycode from users_stream;

-- stop after 4
select name, countrycode from users_stream limit 4;

-- basic aggregate
```

```
select countrycode, count(*) from users_stream group by countrycode;

drop stream if exists users_stream delete topic;

list streams;

show topics;
```

## Create a stream with JSON

---

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-
factor 1 --topic USERPROFILE

kafka-console-producer --broker-list localhost:9092 --topic USERPROFILE << EOF
{"userid": 1000, "firstname":"Alison", "lastname":"Smith", "countrycode":"GB",
"rating":4.7}
EOF

kafka-console-producer --broker-list localhost:9092 --topic USERPROFILE << EOF
{"userid": 1001, "firstname":"Bob", "lastname":"Smith", "countrycode":"US",
"rating":4.2}
EOF
```

At KSQL prompt

```
CREATE STREAM userprofile (userid INT, firstname VARCHAR, lastname VARCHAR,
countrycode VARCHAR, rating DOUBLE) \
  WITH (VALUE_FORMAT = 'JSON', KAFKA_TOPIC = 'USERPROFILE');
```

```
SET 'auto.offset.reset'='earliest';

select firstname, lastname, countrycode, rating from userprofile;

Alison | Smith | GB | 4.7
```

## Manipulate a stream

---

Run a data gen

At UNIX prompt

```
ksql-datagen schema=./datagen/userprofile.avro format=json topic=USERPROFILE
key=userid maxInterval=5000 iterations=100
```

At KSQL prompt

```
-- Review a stream - every 5th row
print 'USERPROFILE' interval 5;
```

## Manipulate a stream

At KSQL prompt

```
ksql> describe userprofile;
```

Name		: USERPROFILE	
Field		Type	
-----			
ROWTIME		BIGINT	(system) <-- NOTE
ROWKEY		VARCHAR(String)	(system) <-- NOTE
USERID		INTEGER	
FIRSTNAME		VARCHAR(String)	
LASTNAME		VARCHAR(String)	
COUNTRYCODE		VARCHAR(String)	
RATING		DOUBLE	

```
select rowtime, firstname from userprofile;
```

- Review *Scalar functions* at <https://docs.confluent.io/current/ksql/docs/developer-guide/syntax-reference.html#scalar-functions>

```
SELECT TIMESTAMPTOSTRING(rowtime, 'dd/MM HH:mm') as createtime, firstname
from userprofile;

select  TIMESTAMPTOSTRING(rowtime, 'dd/MM HH:mm') as createtime, firstname || ' '
|| ucase(lastname) as full_name
from userprofile;
```

## Create a stream from a stream

At KSQL prompt

```
select firstname || ' '
|| ucase( lastname)
|| ' from ' || countrycode
|| ' has a rating of ' || cast(rating as varchar) || ' stars. '
|| case when rating < 2.5 then 'Poor'
        when rating between 2.5 and 4.2 then 'Good'
        else 'Excellent'
end as description
from userprofile;
```

```
Bob FAWCETT from IN has a rating of 4.4 stars. | Excellent
Heidi COEN from US has a rating of 4.9 stars. | Excellent
Bob FAWCETT from IN has a rating of 2.2 stars. | Poor
```

At KSQL prompt

Review the script `user_profile_pretty.ksql`

```
list streams;

run script './user_profile_pretty.ksql';

list streams;

describe extended user_profile_pretty;

select description from user_profile_pretty;

drop stream user_profile_pretty;

terminate query CSAS_USER_PROFILE_PRETTY_0;

drop stream user_profile_pretty;

list streams;

drop stream IF EXISTS user_profile_pretty DELETE TOPIC;
```

## Create a table

---

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-
factor 1 --topic COUNTRY-CSV

kafka-console-producer --broker-list localhost:9092 --topic COUNTRY-CSV --property
```

```
"parse.key=true" --property "key.separator=:" << EOF
AU:AU,Australia
IN:IN,India
GB:GB,UK
US:US,United States
EOF
```

At KSQL prompt

```
CREATE TABLE COUNTRYTABLE (countrycode VARCHAR, countryname VARCHAR) WITH
(KAFKA_TOPIC='COUNTRY-CSV', VALUE_FORMAT='DELIMITED', KEY = 'countrycode');

show tables;

describe COUNTRYTABLE;

describe extended COUNTRYTABLE;

SET 'auto.offset.reset'='earliest';

select countrycode, countryname from countrytable;

-- Note the countryname is "UK"
select countrycode, countryname from countrytable where countrycode='GB' limit 1;

-- This does not exist
select countrycode, countryname from countrytable where countrycode='FR';
```

## Update a table

---

One record updated (UK->United Kingdom), one record added (FR)

At UNIX prompt

```
kafka-console-producer --broker-list localhost:9092 --topic COUNTRY-CSV --property
"parse.key=true" --property "key.separator=:" << EOF
GB:GB,United Kingdom
FR:FR,France
EOF
```

At KSQL prompt

```
select countrycode, countryname from countrytable;

-- Note the countryname has changed to "United Kingdom"
select countrycode, countryname from countrytable where countrycode='GB' limit 1;
```

```
-- And now appears
select countrycode, countryname from countrytable where countrycode='FR';
```

## Join

---

Join user stream to country table

At KSQL prompt

```
select up.firstname, up.lastname, up.countrycode, ct.countryname
from USERPROFILE up
left join COUNTRYTABLE ct on ct.countrycode=up.countrycode;

create stream up_joined as
select up.firstname
|| ' ' || ucase(up.lastname)
|| ' from ' || ct.countryname
|| ' has a rating of ' || cast(rating as varchar) || ' stars.' as description
from USERPROFILE up
left join COUNTRYTABLE ct on ct.countrycode=up.countrycode;

select * from up_joined ;
```

## Data Formats

---

Imagine a *complaints* stream of unhappy customers. Explore the different data formats (CSV, JSON, AVRO)

Column	AVRO Type	KSQL Type
customer_name	string	VARCHAR
complaint_type	string	VARCHAR
trip_cost	float	DOUBLE
new_customer	boolean	BOOLEAN

### CSV Delimited

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-
factor 1 --topic COMPLAINTS_CSV

kafka-console-producer --broker-list localhost:9092 --topic COMPLAINTS_CSV << EOF
```

```
Alice, Late arrival, 43.10, true
EOF
```

At KSQL prompt

```
CREATE STREAM complaints_csv (customer_name VARCHAR, complaint_type VARCHAR,
trip_cost DOUBLE, new_customer BOOLEAN) \
  WITH (VALUE_FORMAT = 'DELIMITED', KAFKA_TOPIC = 'COMPLAINTS_CSV');

select * from complaints_csv;
```

## CSV - experience with bad data

At UNIX prompt

```
kafka-console-producer --broker-list localhost:9092 --topic COMPLAINTS_CSV << EOF
Alice, Bob and Carole, Bad driver, 43.10, true
EOF
```

## JSON

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-
factor 1 --topic COMPLAINTS_JSON

kafka-console-producer --broker-list localhost:9092 --topic COMPLAINTS_JSON << EOF
{"customer_name":"Alice, Bob and Carole", "complaint_type":"Bad driver",
"trip_cost": 22.40, "new_customer": true}
EOF
```

At KSQL prompt

```
CREATE STREAM complaints_json (customer_name VARCHAR, complaint_type VARCHAR,
trip_cost DOUBLE, new_customer BOOLEAN) \
  WITH (VALUE_FORMAT = 'JSON', KAFKA_TOPIC = 'COMPLAINTS_JSON');

select * from complaints_json;
```

## JSON - experience with bad data

At UNIX prompt



```
kafka-console-producer --broker-list localhost:9092 --topic COMPLAINTS_JSON << EOF
{"customer_name":"Bad Data", "complaint_type":"Bad driver", "trip_cost": 22.40,
"new_customer": ShouldBeABoolean}
EOF
```

Review the KSQL Server logs `confluent local log ksql-server`

Now look at the *KSQL Server log*. We can see bad data is noticed; but hidden in a conversion error message

```
at [Source: (byte[]){"customer_name":"Bad Data", "complaint_type":"Bad driver",
"trip_cost": 22.40, "new_customer": ShouldBeABoolean}"; line: 1, column: 105]
Caused by: com.fasterxml.jackson.core.JsonParseException: Unrecognized token
'ShouldBeABoolean': was expecting ('true', 'false' or 'null')
```

## AVRO

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-
factor 1 --topic COMPLAINTS_AVRO

kafka-avro-console-producer --broker-list localhost:9092 --topic COMPLAINTS_AVRO \
--property value.schema='
{
  "type": "record",
  "name": "myrecord",
  "fields": [
    {"name": "customer_name", "type": "string" }
    , {"name": "complaint_type", "type": "string" }
    , {"name": "trip_cost", "type": "float" }
    , {"name": "new_customer", "type": "boolean"}
  ]
}' << EOF
{"customer_name":"Carol", "complaint_type":"Late arrival", "trip_cost": 19.60,
"new_customer": false}
EOF
```

At KSQL prompt

```
-- Note no columns or data type specified
create stream complaints_avro with (kafka_topic='COMPLAINTS_AVRO',
value_format='AVRO');

describe extended complaints_avro;
```

## AVRO - experience with bad data

At UNIX prompt - note bad data is noted at serialization time

```
kafka-avro-console-producer --broker-list localhost:9092 --topic COMPLAINTS_AVRO \
--property value.schema='
{
  "type": "record",
  "name": "myrecord",
  "fields": [
    {"name": "customer_name", "type": "string" }
    , {"name": "complaint_type", "type": "string" }
    , {"name": "trip_cost", "type": "float" }
    , {"name": "new_customer", "type": "boolean"}
  ]
}' << EOF
{"customer_name":"Bad Data", "complaint_type":"Bad driver", "trip_cost": 22.40,
"new_customer": ShouldBeABoolean}
EOF
```

## AVRO Schema Evolution

At UNIX prompt

```
# Optional : strart Confluent Control Center
confluent local start

curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions

kafka-avro-console-producer --broker-list localhost:9092 --topic COMPLAINTS_AVRO \
--property value.schema='
{
  "type": "record",
  "name": "myrecord",
  "fields": [
    {"name": "customer_name", "type": "string" }
    , {"name": "complaint_type", "type": "string" }
    , {"name": "trip_cost", "type": "float" }
    , {"name": "new_customer", "type": "boolean"}
    , {"name": "number_of_rides", "type": "int", "default" : 1}
  ]
}' << EOF
{"customer_name":"Ed", "complaint_type":"Dirty car", "trip_cost": 29.10,
"new_customer": false, "number_of_rides": 22}
EOF

curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions
```

```
curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions/1 |
jq '.'

curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions/2 |
jq '.'
```

At KSQL prompt

```
ksql> describe complaints_avro;

Name                               : COMPLAINTS_AVRO
Field                               | Type
-----
ROWTIME                            | BIGINT          (system)
ROWKEY                             | VARCHAR(STRING) (system)
CUSTOMER_NAME                      | VARCHAR(STRING)
COMPLAINT_TYPE                     | VARCHAR(STRING)
TRIP_COST                          | DOUBLE
NEW_CUSTOMER                       | BOOLEAN
-----

ksql> create stream complaints_avro_v2 with (kafka_topic='COMPLAINTS_AVRO',
value_format='AVRO');

ksql> describe complaints_avro_v2;

Name                               : COMPLAINTS_AVRO_V2
Field                               | Type
-----
ROWTIME                            | BIGINT          (system)
ROWKEY                             | VARCHAR(STRING) (system)
CUSTOMER_NAME                      | VARCHAR(STRING)
COMPLAINT_TYPE                     | VARCHAR(STRING)
TRIP_COST                          | DOUBLE
NEW_CUSTOMER                       | BOOLEAN
NUMBER_OF_RIDES                    | INTEGER          <-- *** NOTE new column
-----
```

## Nested JSON

Imagine we have data like this

```
{
  "city": {
    "name": "Sydney",
```

```
"country": "AU",  
"latitude": -33.8688,  
"longitude": 151.2093  
,  
"description": "light rain",  
"clouds": 92,  
"deg": 26,  
"humidity": 94,  
"pressure": 1025.12,  
"rain": 1.25  
}
```

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-  
factor 1 --topic WEATHERNESTED  
  
cat demo-weather.json | kafka-console-producer --broker-list localhost:9092 --  
topic WEATHERNESTED
```

Extract like this - At KSQL prompt

```
SET 'auto.offset.reset'='earliest';  
  
CREATE STREAM weather  
  (city STRUCT <name VARCHAR, country VARCHAR, latitude DOUBLE, longitude  
DOUBLE>,  
  description VARCHAR,  
  clouds BIGINT,  
  deg BIGINT,  
  humidity BIGINT,  
  pressure DOUBLE,  
  rain DOUBLE)  
WITH (KAFKA_TOPIC='WEATHERNESTED', VALUE_FORMAT='JSON');  
  
SELECT city->name AS city_name, city->country AS city_country, city->latitude as  
latitude, city->longitude as longitude, description, rain from weather;
```

## Build rekeyed table

---

- create a table based on rekeyed **city** field from **weather** stream
- At KSQL prompt

```
create stream weatherraw with (value_format='AVRO') as SELECT city->name AS  
city_name, city->country AS city_country, city->latitude as latitude, city-
```

```
>longitude as longitude, description, rain from weather ;

list streams;
-- note AVRO

describe extended weatherraw;
```

Now notice the *Key field*

```
ksql> describe extended weatherraw;
>

Name           : WEATHERRAW
Type           : STREAM
Key field      :                               <- *** NOTE BLANK ***
Key format     : STRING
Timestamp field : Not set - using <ROWTIME>
Value format   : AVRO
Kafka topic    : WEATHERRAW (partitions: 4, replication: 1)
```

```
create stream weatherrekeyed as select * from weatherraw partition by city_name;

describe extended weatherrekeyed;
```

Now notice the *Key field*

```
ksql> describe extended weatherrekeyed;
>

Name           : WEATHERREKEYED
Type           : STREAM
Key field      : CITY_NAME      <- *** Keyed on city ***
Key format     : STRING
Timestamp field : Not set - using <ROWTIME>
Value format   : AVRO
Kafka topic    : WEATHERREKEYED (partitions: 4, replication: 1)
```

```
create table weathernow with (kafka_topic='WEATHERREKEYED', value_format='AVRO',
key='CITY_NAME');

select * from weathernow;

select * from weathernow where city_name = 'San Diego';
```

Let's make it sunny! At UNIX prompt

```
cat demo-weather-changes.json | kafka-console-producer --broker-list  
localhost:9092 --topic WEATHERNESTED
```

At KSQL prompt

```
select * from weathernow where city_name = 'San Diego';
```

## Repartition

---

*When you use KSQL to join streaming data, you must ensure that your streams and tables are co-partitioned, which means that input records on both sides of the join have the same configuration settings for partitions.*

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 2 --replication-  
factor 1 --topic DRIVER_PROFILE  
  
kafka-console-producer --broker-list localhost:9092 --topic DRIVER_PROFILE << EOF  
{ "driver_name": "Mr. Speedy", "countrycode": "AU", "rating": 2.4 }  
EOF
```

At KSQL prompt

```
CREATE STREAM DRIVER_PROFILE (driver_name VARCHAR, countrycode VARCHAR, rating  
DOUBLE)  
  WITH (VALUE_FORMAT = 'JSON', KAFKA_TOPIC = 'DRIVER_PROFILE');
```

  

```
select dp.driver_name, ct.countryname, dp.rating  
from DRIVER_PROFILE dp  
left join COUNTRYTABLE ct on ct.countrycode=dp.countrycode;
```

  

Can't join DRIVER\_PROFILE with COUNTRYTABLE since the number of partitions don't match. DRIVER\_PROFILE partitions = 2; COUNTRYTABLE partitions = 1. Please repartition either one so that the number of partitions match.

We can fix this by co-partitioning, use the PARTITION BY clause. At KSQL prompt

```
create stream driverprofile_rekeyed with (partitions=1) as select * from
DRIVER_PROFILE partition by driver_name;

select dp2.driver_name, ct.countryname, dp2.rating
from DRIVERPROFILE_REKEYED dp2
left join COUNTRYTABLE ct on ct.countrycode=dp2.countrycode;
```

## Concat Topics with INSERT

---

- create stream of requested rides in Europe using data gen
- create stream of requested rides in USA using data gen
- combine into single stream of all requested rides using *INSERT*

At UNIX prompt

```
ksql-datagen schema=./datagen/riderequest-europe.avro format=avro
topic=riderequest-europe key=rideid maxInterval=5000 iterations=100

ksql-datagen schema=./datagen/riderequest-america.avro format=avro
topic=riderequest-america key=rideid maxInterval=5000 iterations=100
```

At KSQL prompt

```
create stream rr_america_raw with (kafka_topic='riderequest-america',
value_format='avro');

create stream rr_europe_raw with (kafka_topic='riderequest-europe',
value_format='avro');

select * from rr_america_raw;

select * from rr_europe_raw;

create stream rr_world as select 'Europe' as data_source, * from rr_europe_raw;

insert into rr_world      select 'Americas' as data_source, * from rr_america_raw;

select * from rr_world;
```

## Windows

---

- how many requests are arriving each time period

- At KSQL prompt

```
select data_source, city_name, count(*)
from rr_world
window tumbling (size 60 seconds)
group by data_source, city_name;
```

```
select data_source, city_name, COLLECT_LIST(user)
from rr_world
window tumbling (size 60 seconds)
group by data_source, city_name;
```

```
select data_source, city_name, COLLECT_LIST(user)
from rr_world WINDOW SESSION (60 SECONDS)
group by data_source, city_name;
```

```
select TIMESTAMPTOSTRING(WindowStart(), 'HH:mm:ss')
, TIMESTAMPTOSTRING(WindowEnd(), 'HH:mm:ss')
, data_source
, TOPK(city_name, 3)
, count(*)
FROM rr_world
WINDOW TUMBLING (SIZE 1 minute);
```

## Geospacial

---

- create stream - distance of car to waiting rider
- At KSQL prompt

```
select * from rr_world;

describe rr_world;

create stream requested_journey as
select rr.latitude as from_latitude
, rr.longitude as from_longitude
, rr.user
, rr.city_name as city_name
, w.city_country
, w.latitude as to_latitude
, w.longitude as to_longitude
```



```
, w.description as weather_description
, w.rain
from rr_world rr
left join weathernow w on rr.city_name = w.city_name;

create stream ridetodest as
select user
, city_name
, city_country
, weather_description
, rain
, GEO_DISTANCE(from_latitude, from_longitude, to_latitude, to_longitude, 'km') as
dist
from requested_journey;
```

```
select user || ' is travelling ' || cast(round(dist) as varchar) || ' km to ' ||
city_name || ' where the weather is reported as ' || weather_description
from ridetodest;
```

Alice is at (52,0) and is travelling 215 km to Manchester where it is SUNNY  
Heidi is at (51,-1) and is travelling 88 km to London where it is heavy rain  
Grace is at (50,-1) and is travelling 138 km to London where it is heavy rain

## UDF - Build and deploy KSQL User Defined Anomaly Functions

---

- write a UDF to calculate drive time based on
  - distance to travel
  - weather conditions

### Compile Code to Create Anomaly Functions

- Have a look at the file `java/src/main/java/com/vsimon/kafka/streams/TaxiWait.java`
- If you don't want to compile the code; just copy the JAR from `java/pre-compiled/ksql-udf-taxi-1.0.jar`
- Download Maven and follow the installation instructions (<https://maven.apache.org/>)

```
cd java
mvn clean package
ls target/ksql-udf-taxi-1.0.jar
```

### Deploy KSQL User Defined Functions

Find the location of your extension directory. From KSQL

```
ksql> LIST PROPERTIES;
```

Property	Effective Value
-----	
. . .	
ksql.extension.dir	ext <-- *** Look for this
. . .	

```
# Stop (just the) KSQL-Server
confluent local stop ksql-server

# Create an ext (extensions) directory in ${CONFLUENT_HOME}/ext
mkdir /opt/confluent/ext

# build ksql-udf-taxi.jar as above and copy into ext directory
cp target/ksql-udf-taxi-1.0.jar /opt/confluent/ext

# or to use the pre-compile one
cp pre-compiled/ksql-udf-taxi-1.0.jar /opt/confluent/ext

# Restart KSQL server
confluent local start ksql-server
```

## Check KSQL User Defined Functions Available

Start **ksql** client and verify

```
ksql> list functions;
```

Function Name	Type
-----	
. . .	
SUM	AGGREGATE
TAXI_WAIT	SCALAR <--- You need this one
TIMESTAMPToString	SCALAR

```
ksql> DESCRIBE FUNCTION TAXI_WAIT;
```

```
Name      : TAXI_WAIT
Overview  : Return expected wait time in minutes
Type      : scalar
Jar       : /etc/ksql/ext/ksql-udf-taxi-1.0.jar
Variations :

          Variation : TAXI_WAIT(VARCHAR, DOUBLE)
```

```
Returns      : DOUBLE
Description : Given weather and distance return expected wait time in
minutes
```

## Use the UDF

```
describe ridetodest;

select user
, round(dist) as dist
, weather_description
, round(TAXI_WAIT(weather_description, dist)) as taxi_wait_min
from ridetodest;
```

```
select user
|| ' will be waiting ' || cast(round(TAXI_WAIT(weather_description, dist)) as
varchar)
|| ' minutes for their trip of '
|| cast(round(dist) as varchar) || ' km to ' || city_name
|| ' where it is ' || weather_description
from ridetodest;
```

```
Heidi will be waiting 14 minutes for their trip of 358 km to Bristol where it is
light rain
Bob will be waiting 4 minutes for their trip of 218 km to Manchester where it is
SUNNY
Frank will be waiting 15 minutes for their trip of 193 km to London where it is
heavy rain
```